

## ✓ Introduction to Keras for engineers

**Author:** [fchollet](#)

**Date created:** 2023/07/10

**Last modified:** 2023/07/10

**Description:** First contact with Keras 3.

## ✓ Introduction

Keras 3 is a deep learning framework works with TensorFlow, JAX, and PyTorch interchangeably. This notebook will walk you through key Keras 3 workflows.

Let's start by installing Keras 3:

```
1 !pip install keras==3.0.0 --upgrade --quiet
```

## ✓ Setup

We're going to be using the JAX backend here -- but you can edit the string below to "tensorflow" or "torch" and hit "Restart runtime", and the whole notebook will run just the same! This entire guide is backend-agnostic.

```
1 import numpy as np
2 import os
3
4 os.environ["KERAS_BACKEND"] = "jax"
5
6 # Note that Keras should only be imported after the backend
7 # has been configured. The backend cannot be changed once the
8 # package is imported.
9 import keras
```

## ✓ A first example: A MNIST convnet

Let's start with the Hello World of ML: training a convnet to classify MNIST digits.

Here's the data:

```
1 # Load the data and split it between train and test sets
2 (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
3
4 # Scale images to the [0, 1] range
5 x_train = x_train.astype("float32") / 255
6 x_test = x_test.astype("float32") / 255
7 # Make sure images have shape (28, 28, 1)
8 x_train = np.expand_dims(x_train, -1)
9 x_test = np.expand_dims(x_test, -1)
10 print("x_train shape:", x_train.shape)
11 print("y_train shape:", y_train.shape)
12 print(x_train.shape[0], "train samples")
13 print(x_test.shape[0], "test samples")
```

Here's our model.

Different model-building options that Keras offers include:

- [The Sequential API](#) (what we use below)
- [The Functional API](#) (most typical)
- [Writing your own models yourself via subclassing](#) (for advanced use cases)

```
1 # Model parameters
2 num_classes = 10
3 input_shape = (28, 28, 1)
4
5 model = keras.Sequential(
6     [
7         keras.layers.Input(shape=input_shape),
8         keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
9         keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
10        keras.layers.MaxPooling2D(pool_size=(2, 2)),
11        keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
12        keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
13        keras.layers.GlobalAveragePooling2D(),
14        keras.layers.Dropout(0.5),
15        keras.layers.Dense(num_classes, activation="softmax"),
16    ]
17 )
```

Here's our model summary:

```
1 model.summary()
```

We use the `compile()` method to specify the optimizer, loss function, and the metrics to monitor. Note that with the JAX and TensorFlow backends, XLA compilation is turned on by default.

```
1 model.compile(
2     loss=keras.losses.SparseCategoricalCrossentropy(),
3     optimizer=keras.optimizers.Adam(learning_rate=1e-3),
4     metrics=[
5         keras.metrics.SparseCategoricalAccuracy(name="acc"),
6     ],
7 )
```

Let's train and evaluate the model. We'll set aside a validation split of 15% of the data during training to monitor generalization on unseen data.

```
1 batch_size = 128
2 epochs = 20
3
4 callbacks = [
5     keras.callbacks.ModelCheckpoint(filepath="model_at_epoch_{epoch}.keras"),
6     keras.callbacks.EarlyStopping(monitor="val_loss", patience=2),
7 ]
8
9 model.fit(
10     x_train,
11     y_train,
12     batch_size=batch_size,
13     epochs=epochs,
14     validation_split=0.15,
15     callbacks=callbacks,
16 )
17 score = model.evaluate(x_test, y_test, verbose=0)
```

During training, we were saving a model at the end of each epoch. You can also save the model in its latest state like this:

```
1 model.save("final_model.keras")
```

And reload it like this:

```
1 model = keras.saving.load_model("final_model.keras")
```

Next, you can query predictions of class probabilities with `predict()`:

```
1 predictions = model.predict(x_test)
```

That's it for the basics!

## ✓ Writing cross-framework custom components

Keras enables you to write custom Layers, Models, Metrics, Losses, and Optimizers that work across TensorFlow, JAX, and PyTorch with the same codebase. Let's take a look at custom layers first.

The `keras.ops` namespace contains:

- An implementation of the NumPy API, e.g. `keras.ops.stack` or `keras.ops.matmul`.
- A set of neural network specific ops that are absent from NumPy, such as `keras.ops.conv` or `keras.ops.binary_crossentropy`.

Let's make a custom `Dense` layer that works with all backends:

```
1
2 class MyDense(keras.layers.Layer):
3     def __init__(self, units, activation=None, name=None):
4         super().__init__(name=name)
5         self.units = units
6         self.activation = keras.activations.get(activation)
7
8     def build(self, input_shape):
9         input_dim = input_shape[-1]
10        self.w = self.add_weight(
11            shape=(input_dim, self.units),
12            initializer=keras.initializers.GlorotNormal(),
13            name="kernel",
14            trainable=True,
15        )
16
17        self.b = self.add_weight(
18            shape=(self.units,),
19            initializer=keras.initializers.Zeros(),
20            name="bias",
21            trainable=True,
22        )
23
24    def call(self, inputs):
25        # Use Keras ops to create backend-agnostic layers/metrics/etc.
26        x = keras.ops.matmul(inputs, self.w) + self.b
27        return self.activation(x)
28
```

Next, let's make a custom `Dropout` layer that relies on the `keras.random` namespace:

```
1
2 class MyDropout(keras.layers.Layer):
3     def __init__(self, rate, name=None):
4         super().__init__(name=name)
5         self.rate = rate
6         # Use seed_generator for managing RNG state.
7         # It is a state element and its seed variable is
8         # tracked as part of `layer.variables`.
9         self.seed_generator = keras.random.SeedGenerator(1337)
10
11     def call(self, inputs):
12         # Use `keras.random` for random ops.
13         return keras.random.dropout(inputs, self.rate, seed=self.seed_generator)
14
```

Next, let's write a custom subclassed model that uses our two custom layers:

```
1
2 class MyModel(keras.Model):
3     def __init__(self, num_classes):
4         super().__init__()
5         self.conv_base = keras.Sequential(
6             [
7                 keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
8                 keras.layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
9                 keras.layers.MaxPooling2D(pool_size=(2, 2)),
10                keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
11                keras.layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
12                keras.layers.GlobalAveragePooling2D(),
13            ]
14        )
15        self.dp = MyDropout(0.5)
16        self.dense = MyDense(num_classes, activation="softmax")
17
18    def call(self, x):
19        x = self.conv_base(x)
20        x = self.dp(x)
21        return self.dense(x)
22
```

Let's compile it and fit it:

```
1 model = MyModel(num_classes=10)
2 model.compile(
3     loss=keras.losses.SparseCategoricalCrossentropy(),
4     optimizer=keras.optimizers.Adam(learning_rate=1e-3),
5     metrics=[
6         keras.metrics.SparseCategoricalAccuracy(name="acc"),
7     ],
8 )
9
10 model.fit(
11     x_train,
12     y_train,
13     batch_size=batch_size,
14     epochs=1, # For speed
15     validation_split=0.15,
16 )
```

## ✓ Training models on arbitrary data sources

All Keras models can be trained and evaluated on a wide variety of data sources, independently of the backend you're using. This includes:

- NumPy arrays
- Pandas dataframes
- TensorFlow `tf.data.Dataset` objects
- PyTorch `DataLoader` objects
- Keras `PyDataset` objects

They all work whether you're using TensorFlow, JAX, or PyTorch as your Keras backend.

Let's try it out with PyTorch `DataLoaders` :

```
1 import torch
2
3 # Create a TensorDataset
4 train_torch_dataset = torch.utils.data.TensorDataset(
5     torch.from_numpy(x_train), torch.from_numpy(y_train)
6 )
7 val_torch_dataset = torch.utils.data.TensorDataset(
8     torch.from_numpy(x_test), torch.from_numpy(y_test)
9 )
10
11 # Create a DataLoader
12 train_dataloader = torch.utils.data.DataLoader(
13     train_torch_dataset, batch_size=batch_size, shuffle=True
14 )
15 val_dataloader = torch.utils.data.DataLoader(
16     val_torch_dataset, batch_size=batch_size, shuffle=False
17 )
18
19 model = MyModel(num_classes=10)
20 model.compile(
21     loss=keras.losses.SparseCategoricalCrossentropy(),
22     optimizer=keras.optimizers.Adam(learning_rate=1e-3),
23     metrics=[
24         keras.metrics.SparseCategoricalAccuracy(name="acc"),
25     ],
26 )
27 model.fit(train_dataloader, epochs=1, validation_data=val_dataloader)
28
```

Now let's try this out with `tf.data`:



```
1 import tensorflow as tf
2
3 train_dataset = (
4     tf.data.Dataset.from_tensor_slices((x_train, y_train))
5     .batch(batch_size)
6     .prefetch(tf.data.AUTOTUNE)
7 )
8 test_dataset = (
9     tf.data.Dataset.from_tensor_slices((x_test, y_test))
10    .batch(batch_size)
11    .prefetch(tf.data.AUTOTUNE)
12 )
13
14 model = MyModel(num_classes=10)
15 model.compile(
16     loss=keras.losses.SparseCategoricalCrossentropy(),
17     optimizer=keras.optimizers.Adam(learning_rate=1e-3),
18     metrics=[
19         keras.metrics.SparseCategoricalAccuracy(name="acc"),
20     ],
21 )
22 model.fit(train_dataset, epochs=1, validation_data=test_dataset)
```

## Further reading

This concludes our short overview of the new multi-backend capabilities of Keras 3. Next, you can learn about:

### How to customize what happens in `fit()`

Want to implement a non-standard training algorithm yourself but still want to benefit from the power and usability of `fit()` ? It's easy to customize `fit()` to support arbitrary use cases:

- [Customizing what happens in `fit\(\)` with TensorFlow](#)
- [Customizing what happens in `fit\(\)` with JAX](#)
- [Customizing what happens in `fit\(\)` with PyTorch](#)

### How to write custom training loops

- [Writing a training loop from scratch in TensorFlow](#)

- [Writing a training loop from scratch in JAX](#)
- [Writing a training loop from scratch in PyTorch](#)

